

OpenShift Developer

Architecture Workshop

Introduction to OpenShift for Developers

Usecase description & Quarkus demo-app



[linkedin.com/company/red-hat](https://www.linkedin.com/company/red-hat)



[facebook.com/redhatinc](https://www.facebook.com/redhatinc)



[youtube.com/user/RedHatVideos](https://www.youtube.com/user/RedHatVideos)



twitter.com/RedHat

Self introduction

Name: Wanja Pernath

Email: wpernath@redhat.com

Base: Germany (very close to the Alps)

Role: EMEA Technical Partner Development

Manager - OpenShift and MW

Experience: Years of Consulting, Training,

PreSales at Red Hat and before



Agenda

Agenda

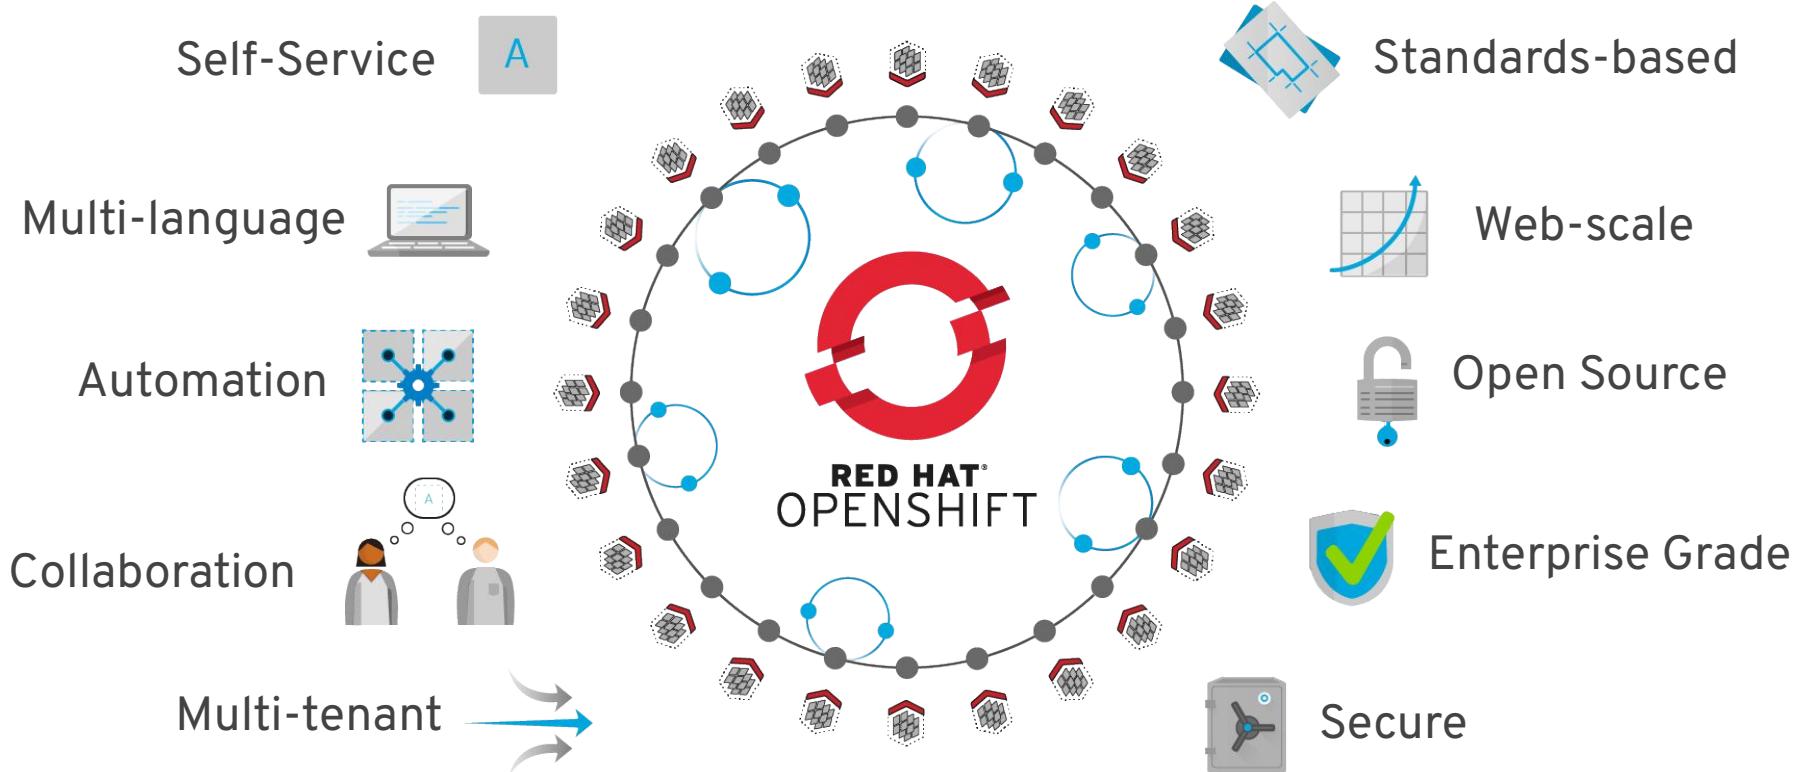
- What is this all about
- The OpenShift Ecosystem for Developers
 - Basics
 - Advanced topics
- How to demo OpenShift for Developers
- Quarkus Business Value
- Quarkus Technical Value

What is this all about?

Learning Goals

- Getting an overview of the whole OpenShift ecosystem for developers
- Learn about the benefits of OpenShift prior plain Kubernetes
- Learn how to effectively and efficiently demo OpenShift for developers
- Learn to use OpenShift and its ecosystem
- Learn to start coding with OpenShift in your preferred language
- Learn to do proper release management with OpenShift; understand the basics and how to include it into your release management process but also learn how to benefit from OpenShift Pipelines and when to use what
- Understanding and using Operators for releases
- Understand and make use of advanced OpenShift features like Istio and Serverless

OpenShift Developers Basics



Value of OpenShift

Monitoring, Logging,
Registry, Router, Telemetry

Cluster Services

Service Mesh, Serverless,
Middleware/Runtimes, ISVs

Application Services

Dev Tools, CI/CD,
Automated Builds, IDE

Developer Services

Automated Operations

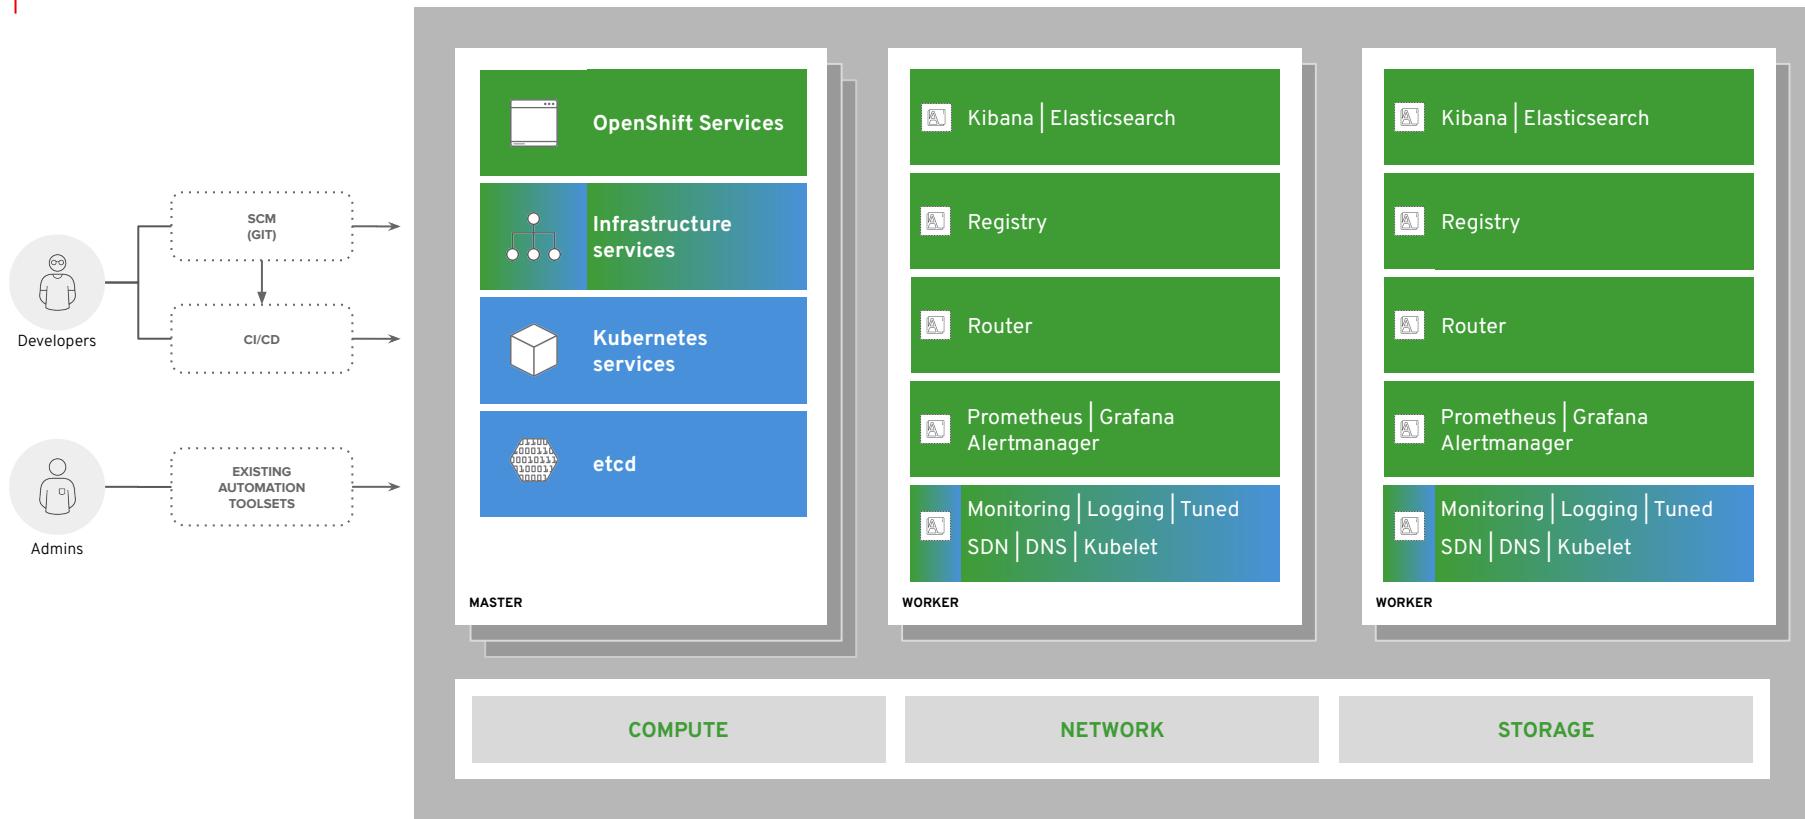
Kubernetes

Red Hat Enterprise Linux | RHEL CoreOS

Best IT Ops Experience

CaaS \longleftrightarrow PaaS \longleftrightarrow FaaS

Best Developer Experience



Overwhelmed? Please see the CNCF Trail Map. That and the interactive landscape are at l.cncf.io

<p>The grid displays a comprehensive overview of CNCF projects across various categories. Projects are represented by icons, with some being 'Graduated' (blue background) and others 'Incubating' (white background). Categories include Database, Streaming & Messaging, Application Definition & Image Build, Continuous Integration & Delivery, Platform, Observability and Analysis, Scheduling & Orchestration, Coordination & Service Discovery, Remote Procedure Call, Service Proxy, API Gateway, Service Mesh, Orchestration & Management, Cloud-Native Storage, Container Runtime, Cloud-Native Network, Runtime, Automation & Configuration, Container Registry, Security & Compliance, Key Management, and Provisioning.</p>											
<p>This section highlights Kubernetes Certified Service Providers, showing a grid of logos for various companies that offer services based on Kubernetes technology.</p>											
<p>This section highlights Kubernetes Training Partners, showing a grid of logos for companies that provide training for Kubernetes.</p>											
<p>This section highlights public cloud providers and infrastructure as a service (IaaS) offerings.</p>											
<p>This section highlights special partners and organizations involved in the CNCF ecosystem.</p>											
<p>This section highlights cloud native technologies and providers.</p>											
<p>This landscape is intended as a map through the previously uncharted terrain of cloud native technologies. There are many routes to deploying a cloud native application, with CNCF Projects representing a particularly well-traveled path.</p> <p>CLOUD NATIVE COMPUTING FOUNDATION</p> <p>CLOUD NATIVE Landscape</p> <p>l.cncf.io</p>											

Core Concepts

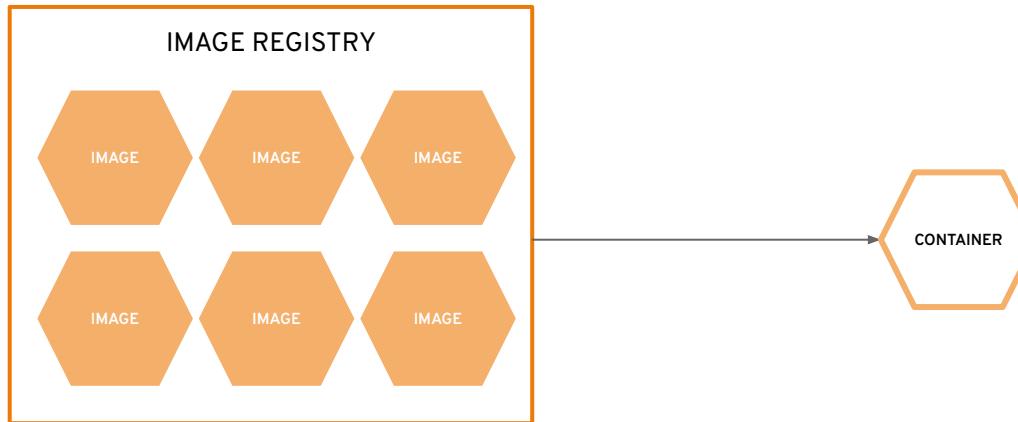
a container is the smallest compute unit



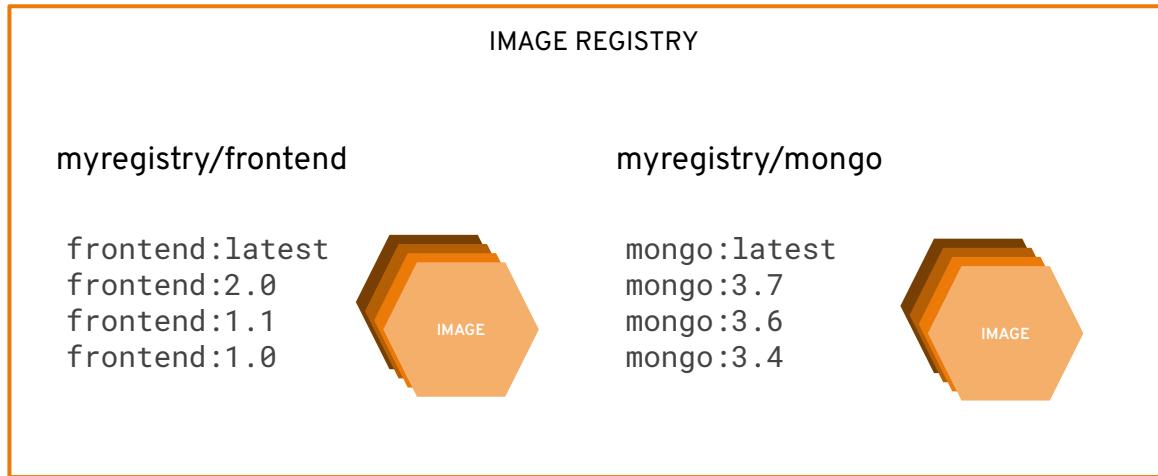
containers are created from container images



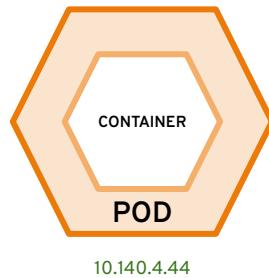
container images are stored in an image registry



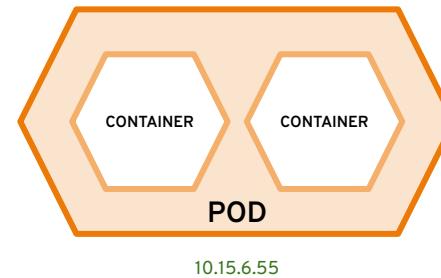
an image repository contains all versions of an image in the image registry



containers are wrapped in pods which are units of deployment and management

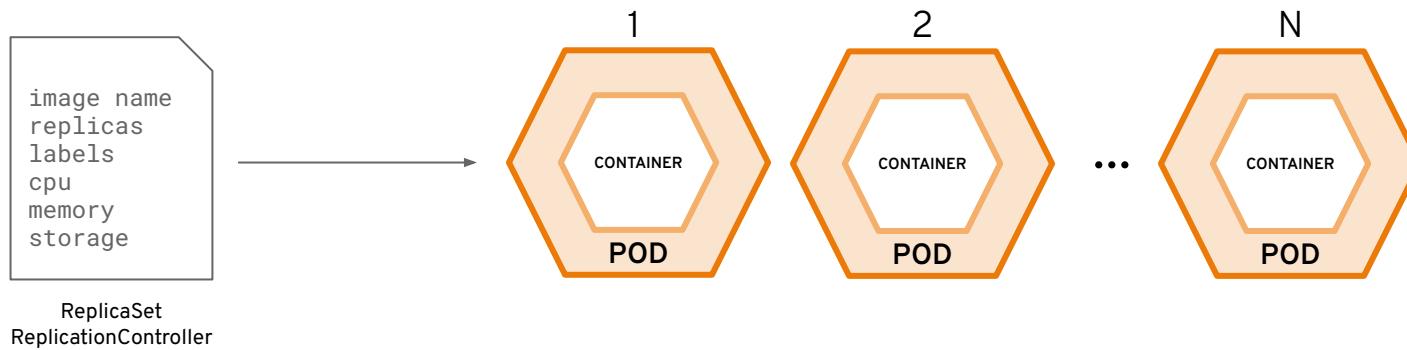


10.140.4.44

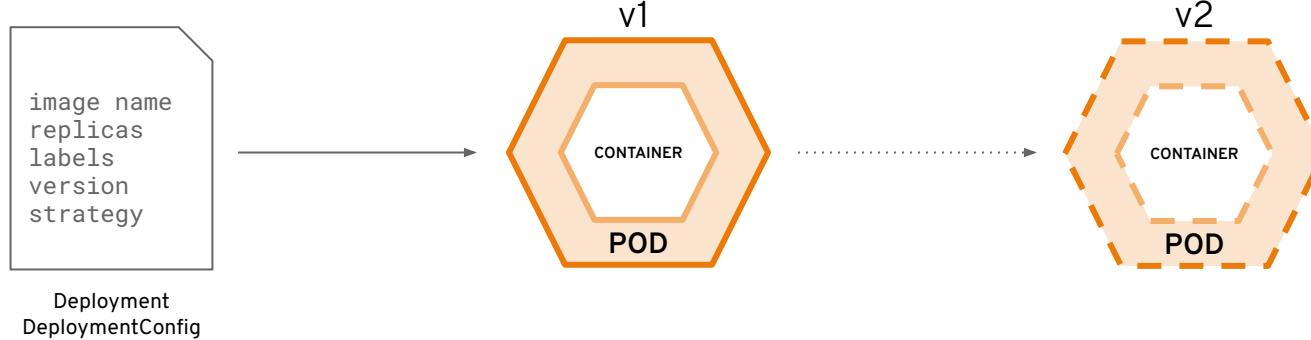


10.15.6.55

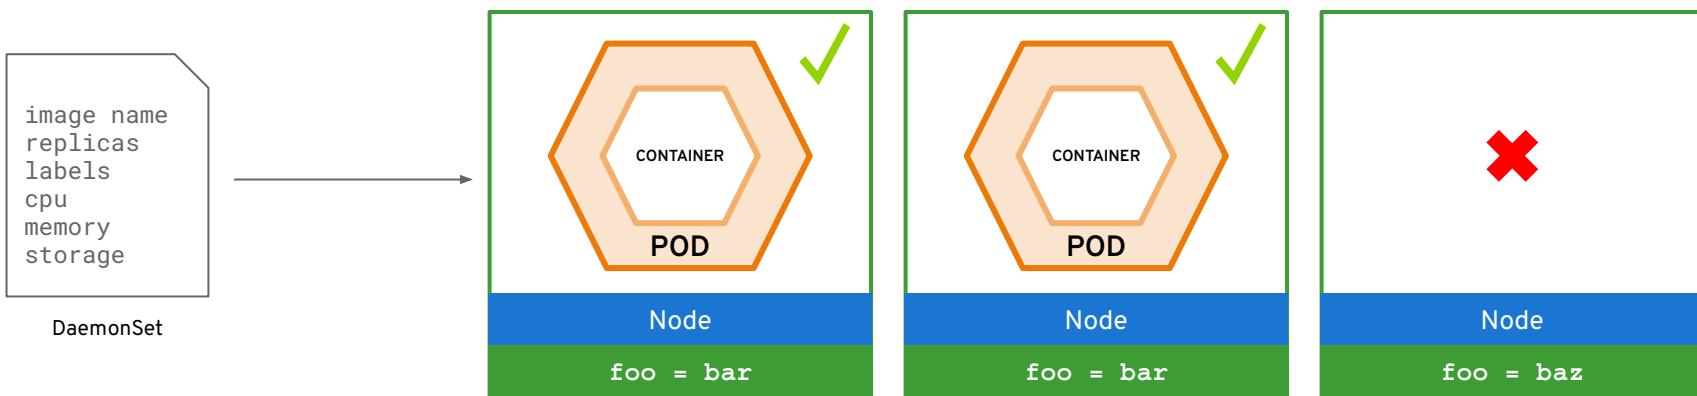
ReplicationControllers & ReplicaSets ensure a specified number of pods are running at any given time



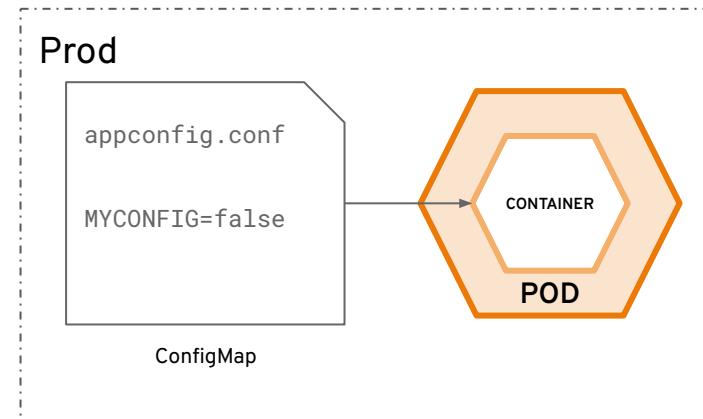
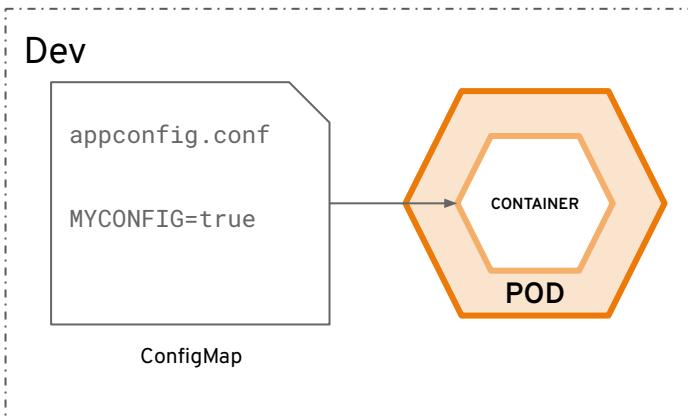
Deployments and DeploymentConfigurations define how to roll out new versions of Pods



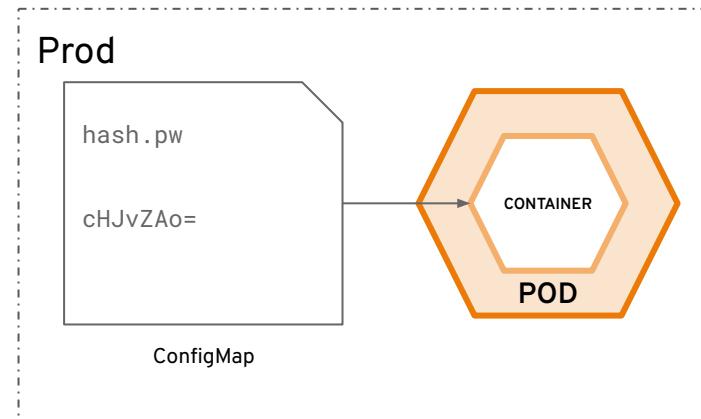
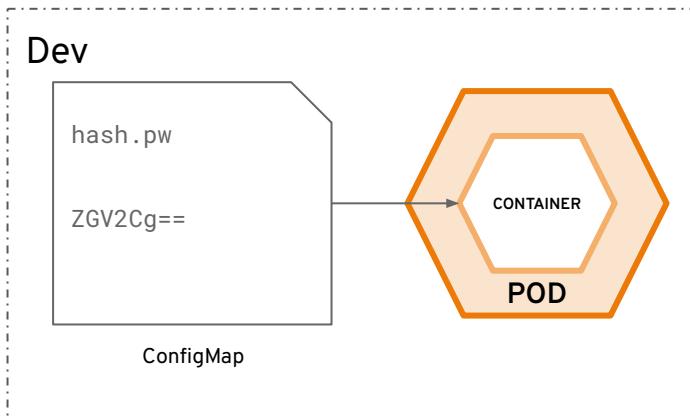
a daemonset ensures that all (or some) nodes run a copy of a pod



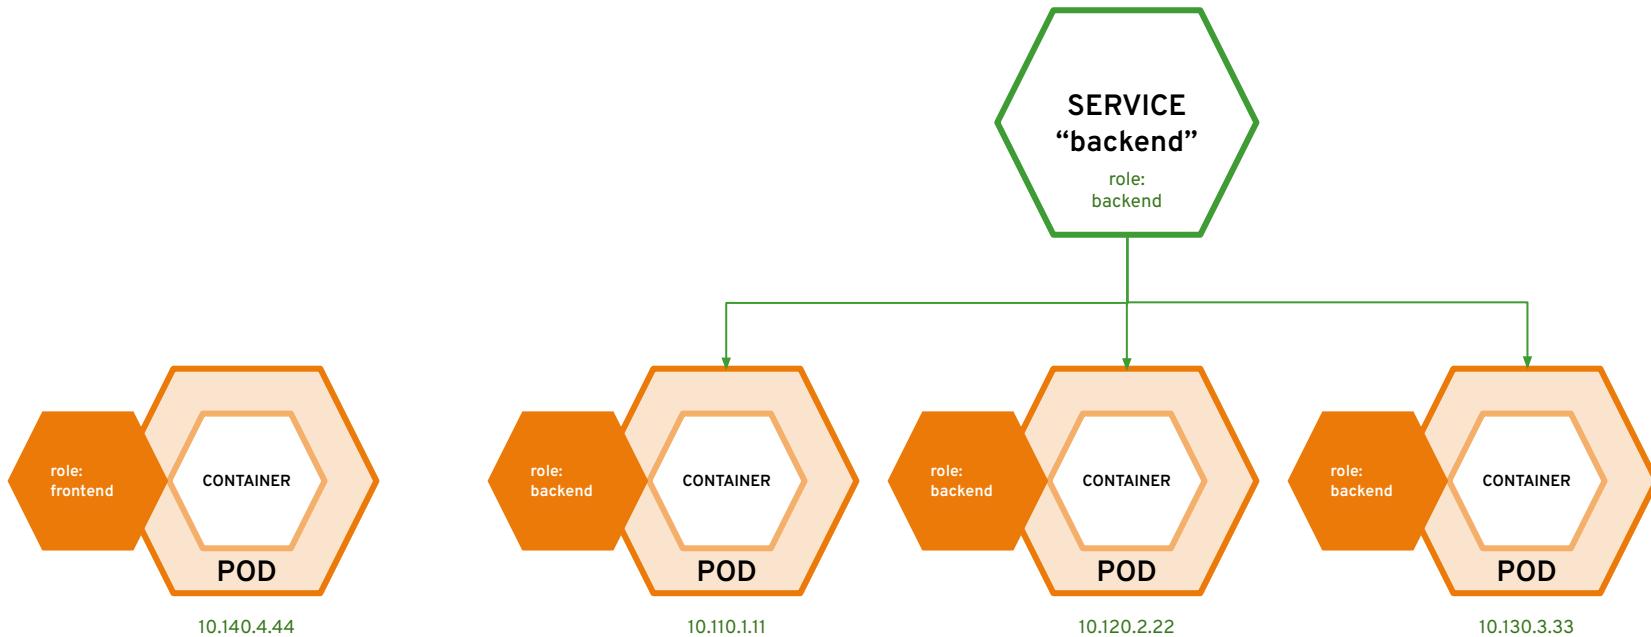
configmaps allow you to decouple configuration artifacts from image content



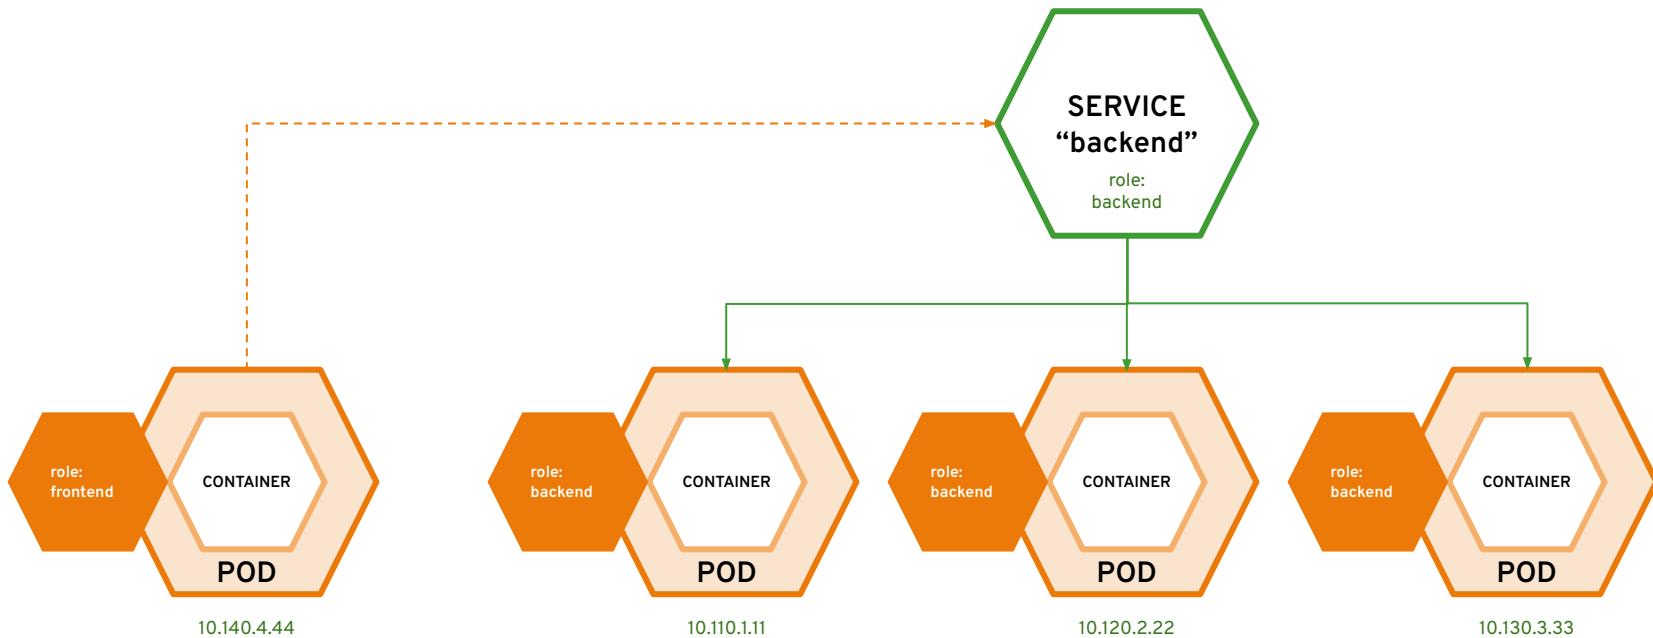
secrets provide a mechanism to hold sensitive information such as passwords



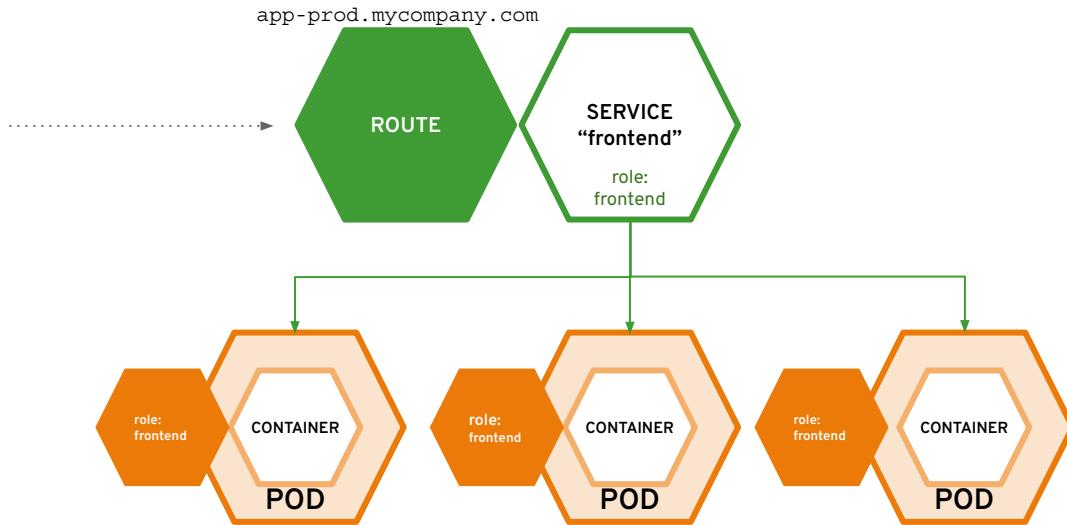
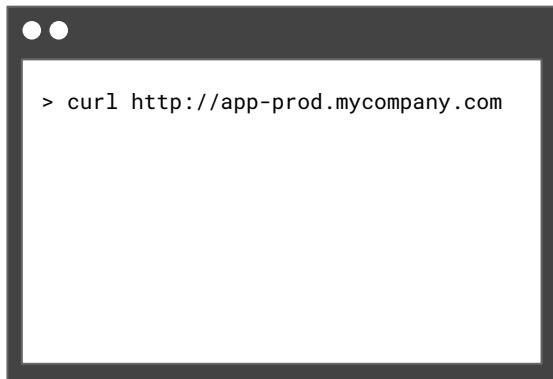
services provide internal load-balancing and service discovery across pods



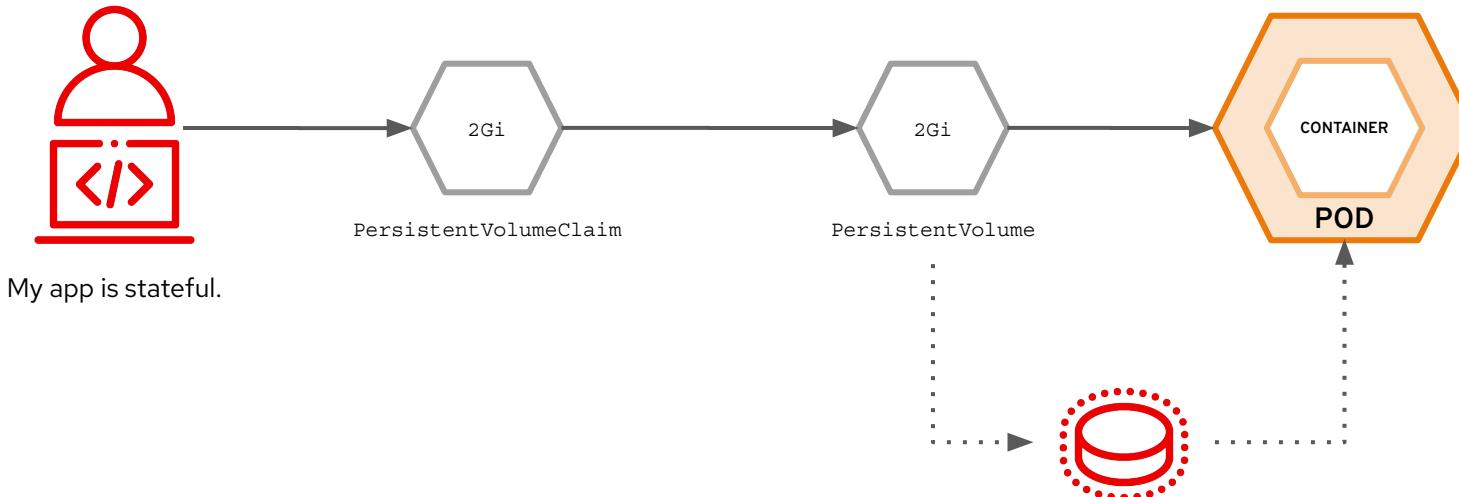
apps can talk to each other via services



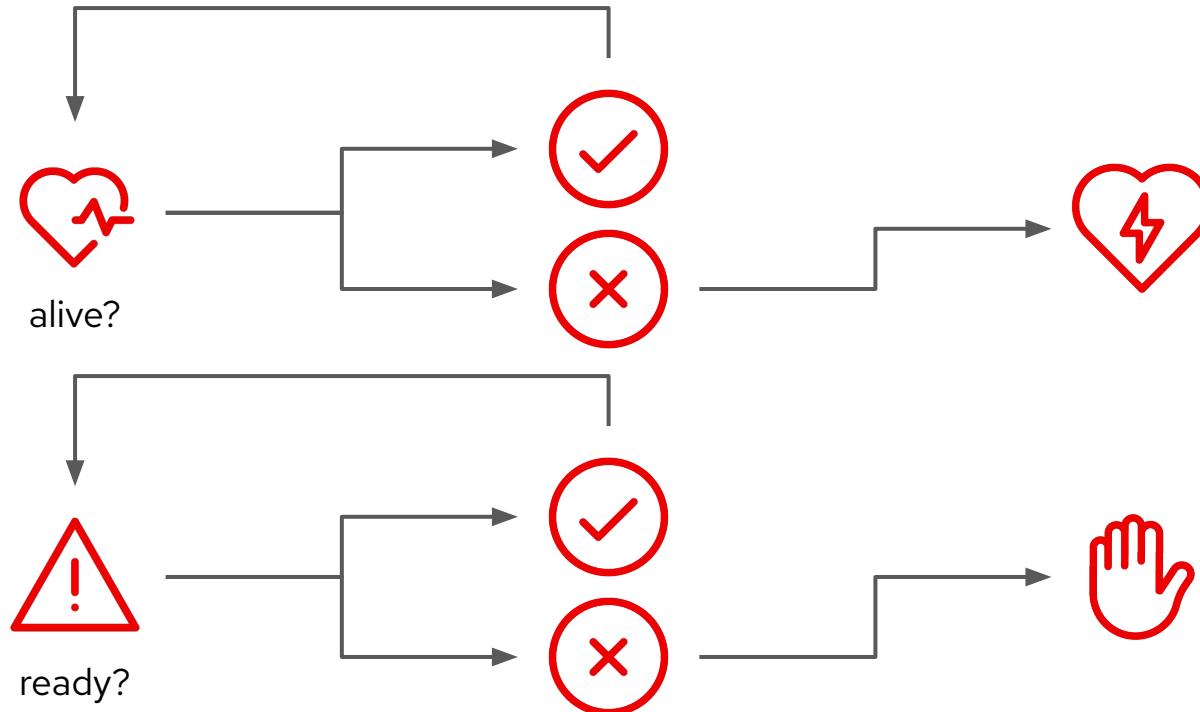
routes make services accessible to clients outside the environment via real-world urls



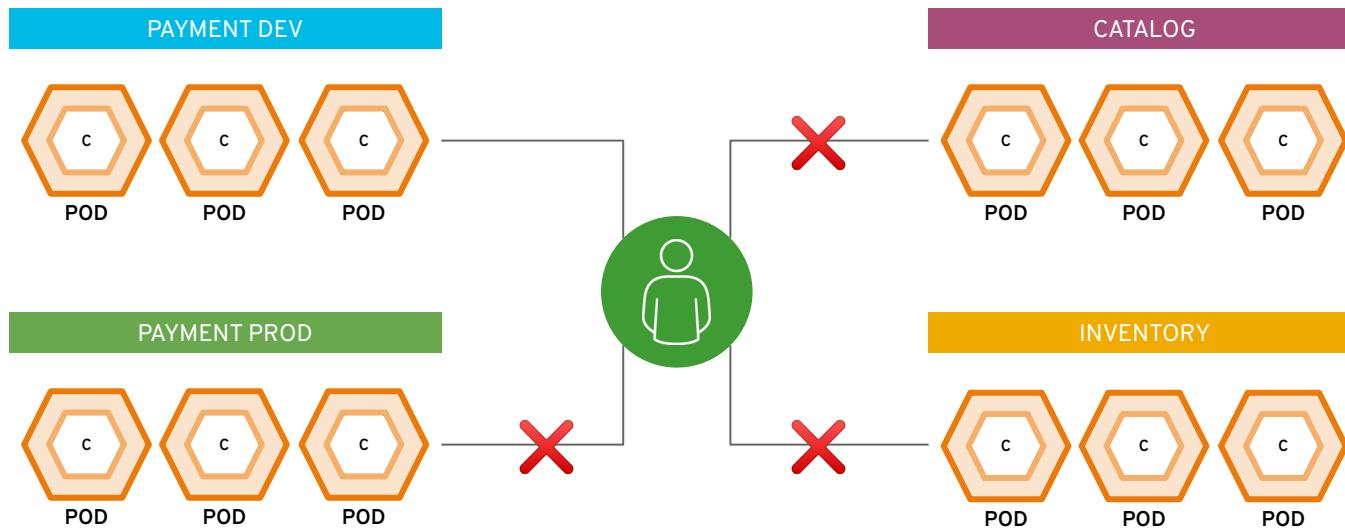
Persistent Volume and Claims



Liveness and Readiness

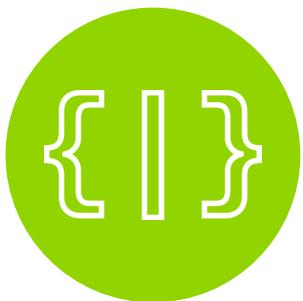


projects isolate apps across environments,
teams, groups and departments



BUILD & DEPLOY

BUILD AND DEPLOY CONTAINER IMAGES



DEPLOY YOUR
SOURCE CODE

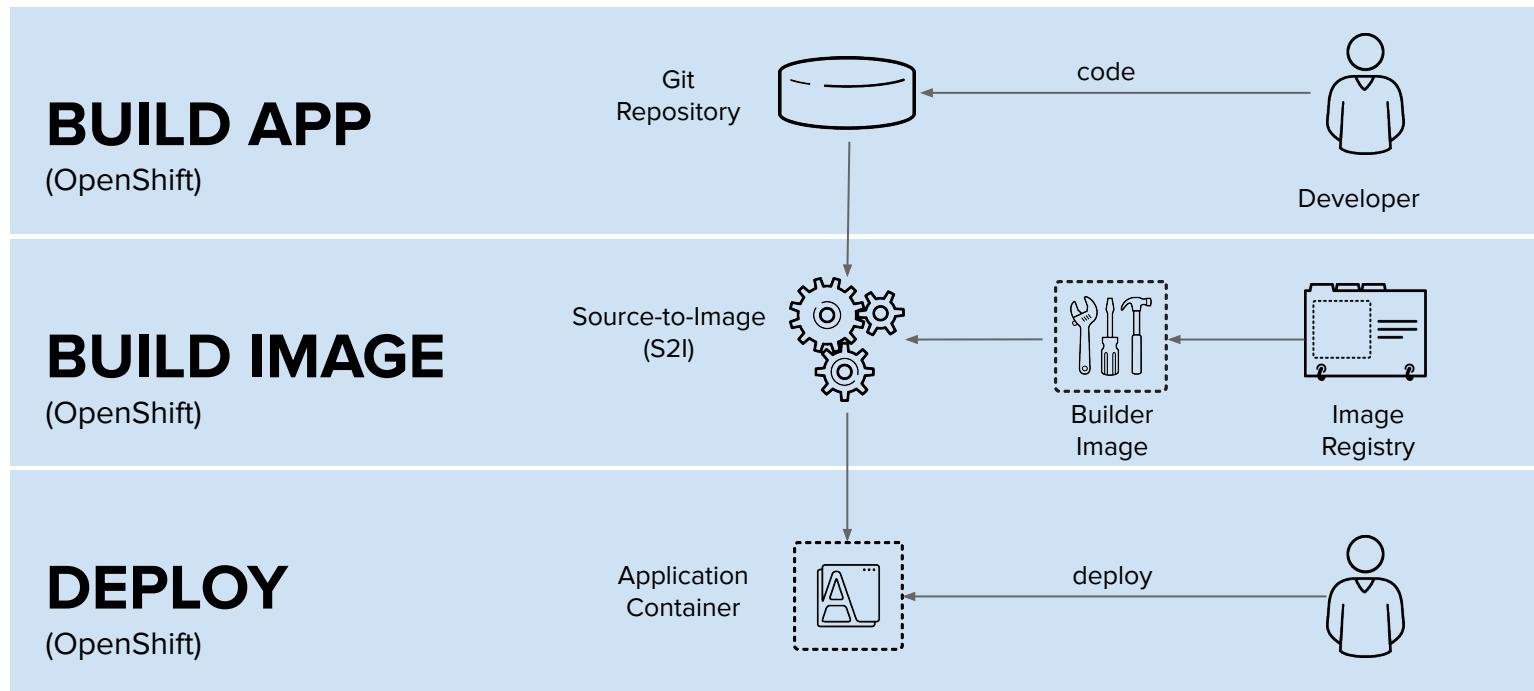


DEPLOY YOUR
APP BINARY

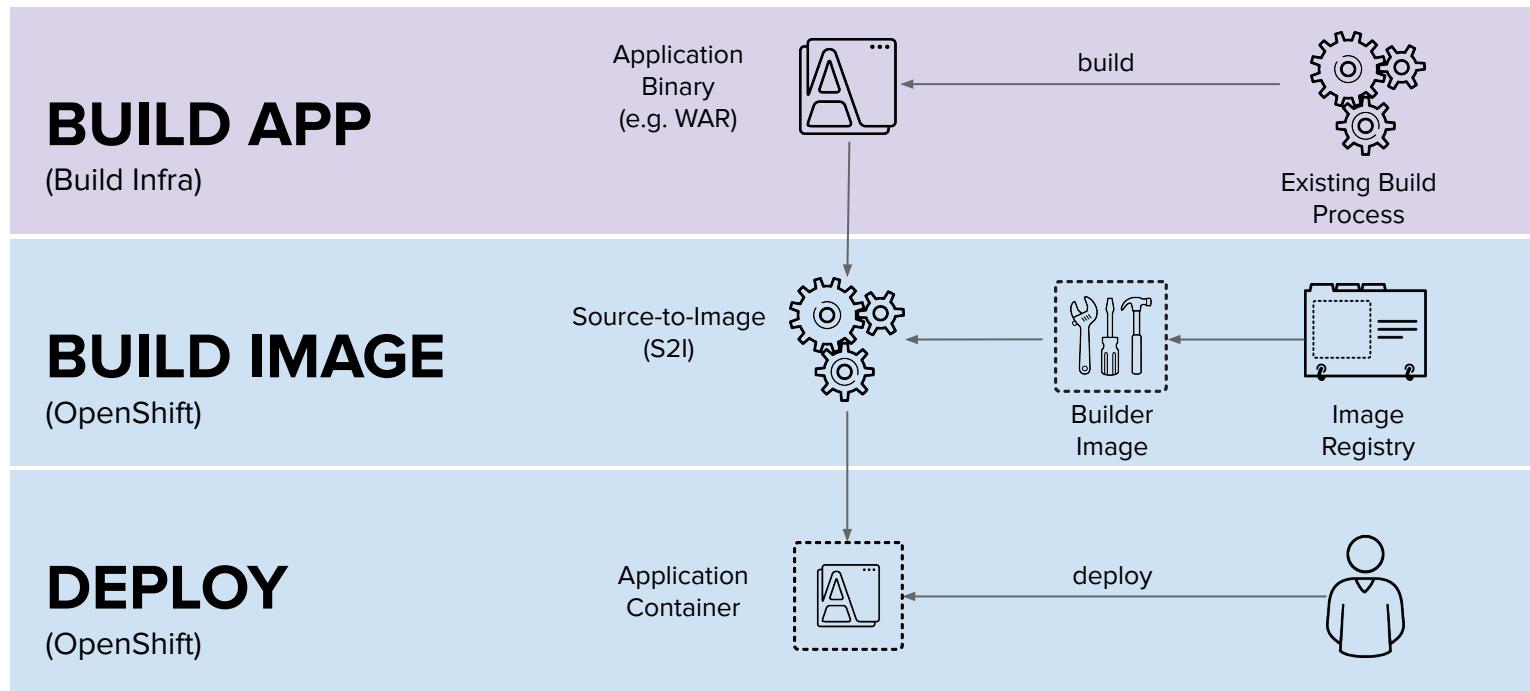


DEPLOY YOUR
CONTAINER IMAGE

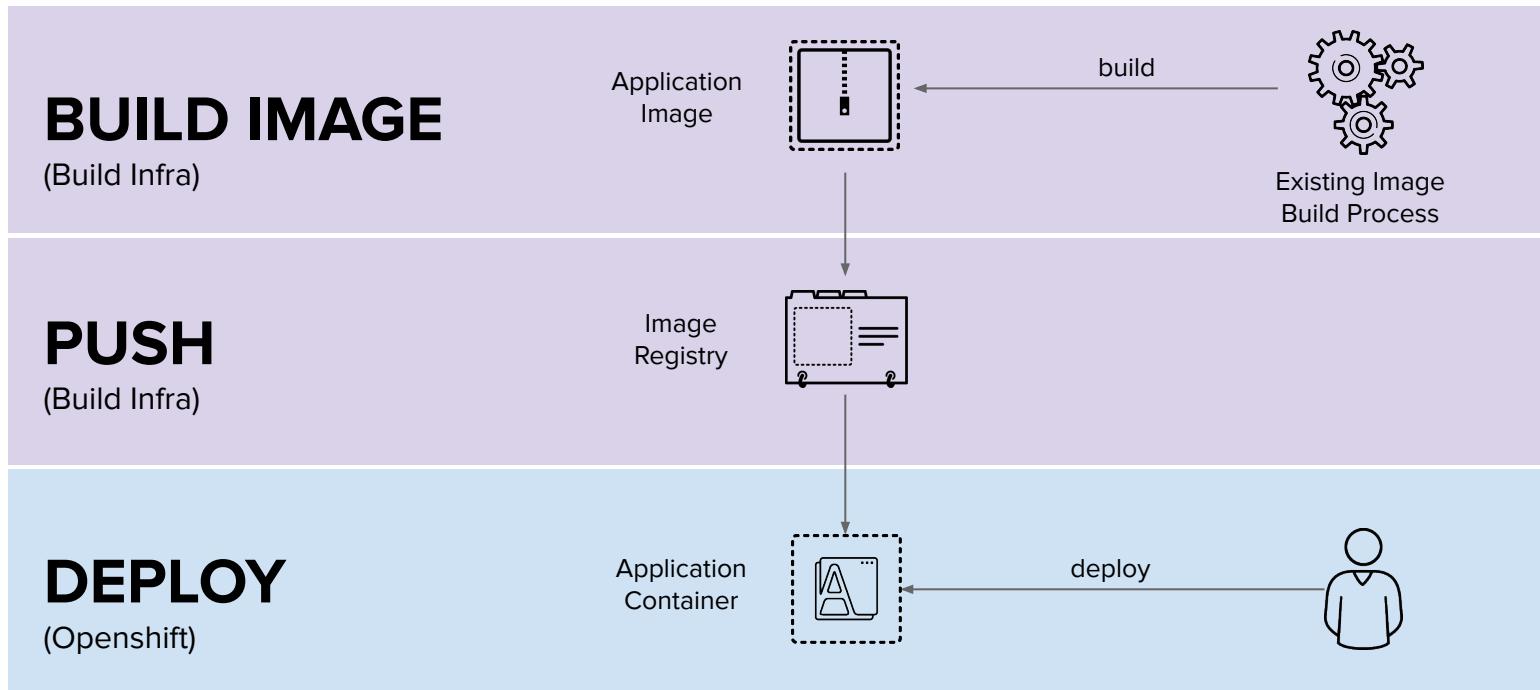
DEPLOY SOURCE CODE WITH SOURCE-TO-IMAGE (S2I)



DEPLOY APP BINARY WITH SOURCE-TO-IMAGE (S2I)

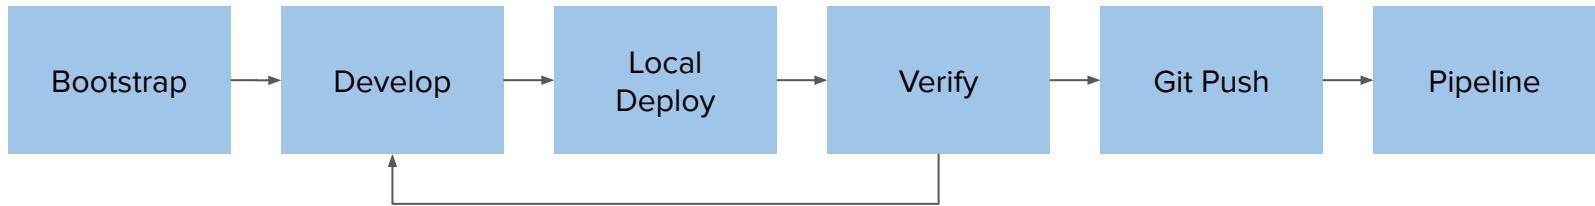


DEPLOY DOCKER IMAGE

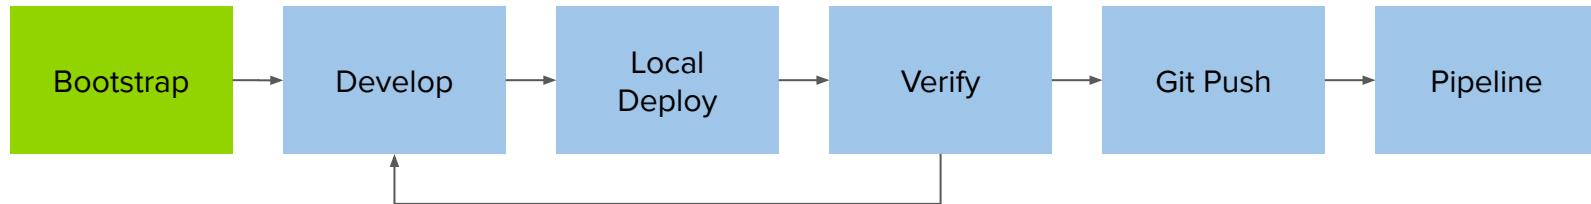


DEVELOPER WORKFLOW

LOCAL DEVELOPMENT WORKFLOW



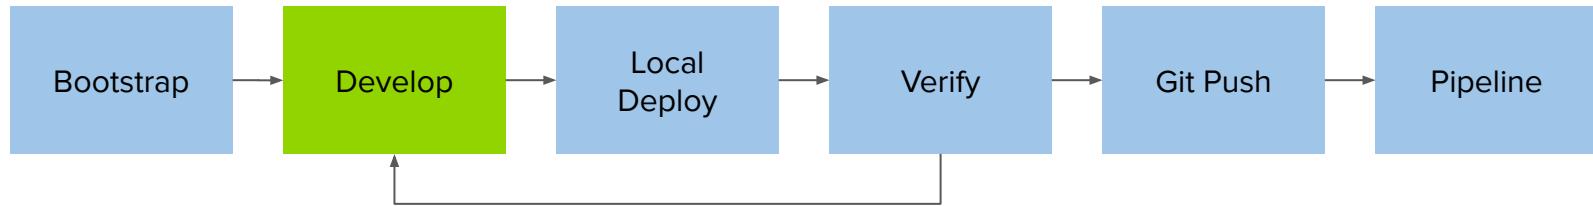
LOCAL DEVELOPMENT WORKFLOW



BOOTSTRAP

- Pick your programming language and application runtime of choice
- Create the project skeleton from scratch or use a generator such as
 - Maven archetypes
 - Quickstarts and Templates
 - OpenShift Generator
 - Spring Initializr

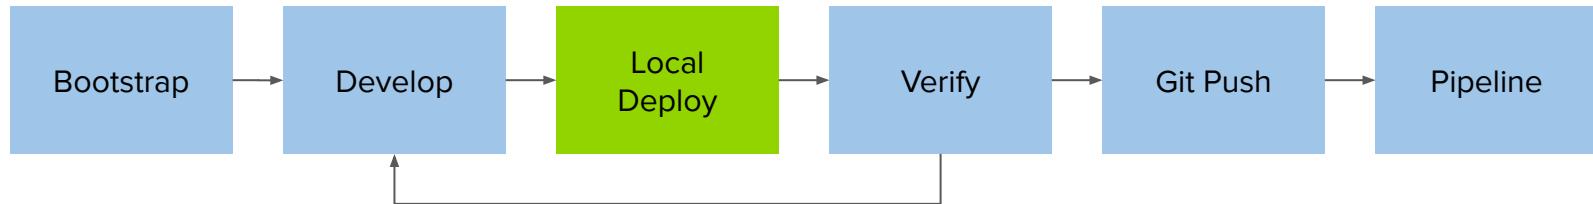
LOCAL DEVELOPMENT WORKFLOW



DEVELOP

- Pick your framework of choice such as Java EE, Spring, Ruby on Rails, Django, Express, ...
- Develop your application code using your editor or IDE of choice
- Build and test your application code locally using your build tools
- Create or generate OpenShift templates or Kubernetes objects

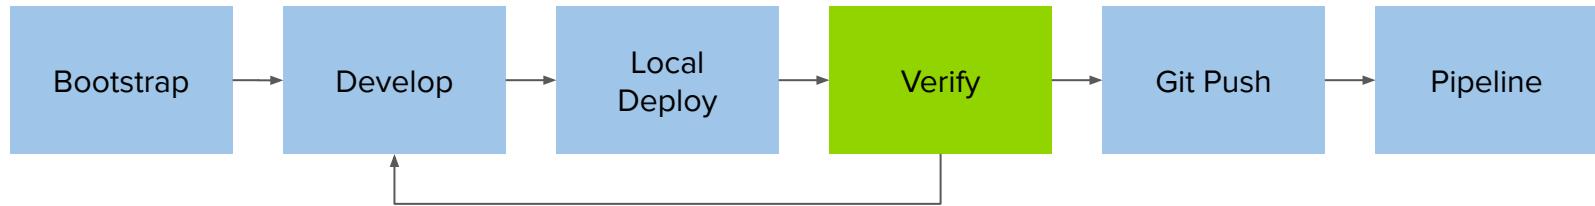
LOCAL DEVELOPMENT WORKFLOW



LOCAL DEPLOY

- Deploy your code on a local OpenShift cluster
 - Red Hat Container Development Kit (CDK), minishift and oc cluster
- Red Hat CDK provides a standard RHEL-based development environment
- Use binary deploy, maven or CLI rsync to push code or app binary directly into containers

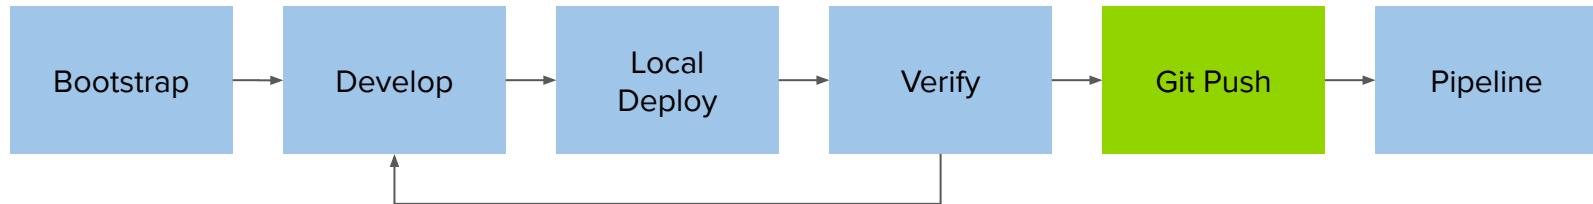
LOCAL DEVELOPMENT WORKFLOW



VERIFY

- Verify your code is working as expected
- Run any type of tests that are required with or without other components (database, etc)
- Based on the test results, change code, deploy, verify and repeat

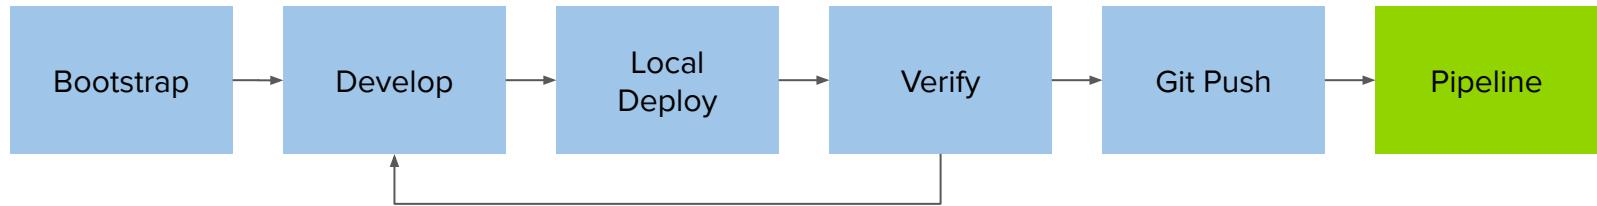
LOCAL DEVELOPMENT WORKFLOW



GIT PUSH

- Push the code and configuration to the Git repository
- If using Fork & Pull Request workflow, create a Pull Request
- If using code review workflow, participate in code review discussions

LOCAL DEVELOPMENT WORKFLOW

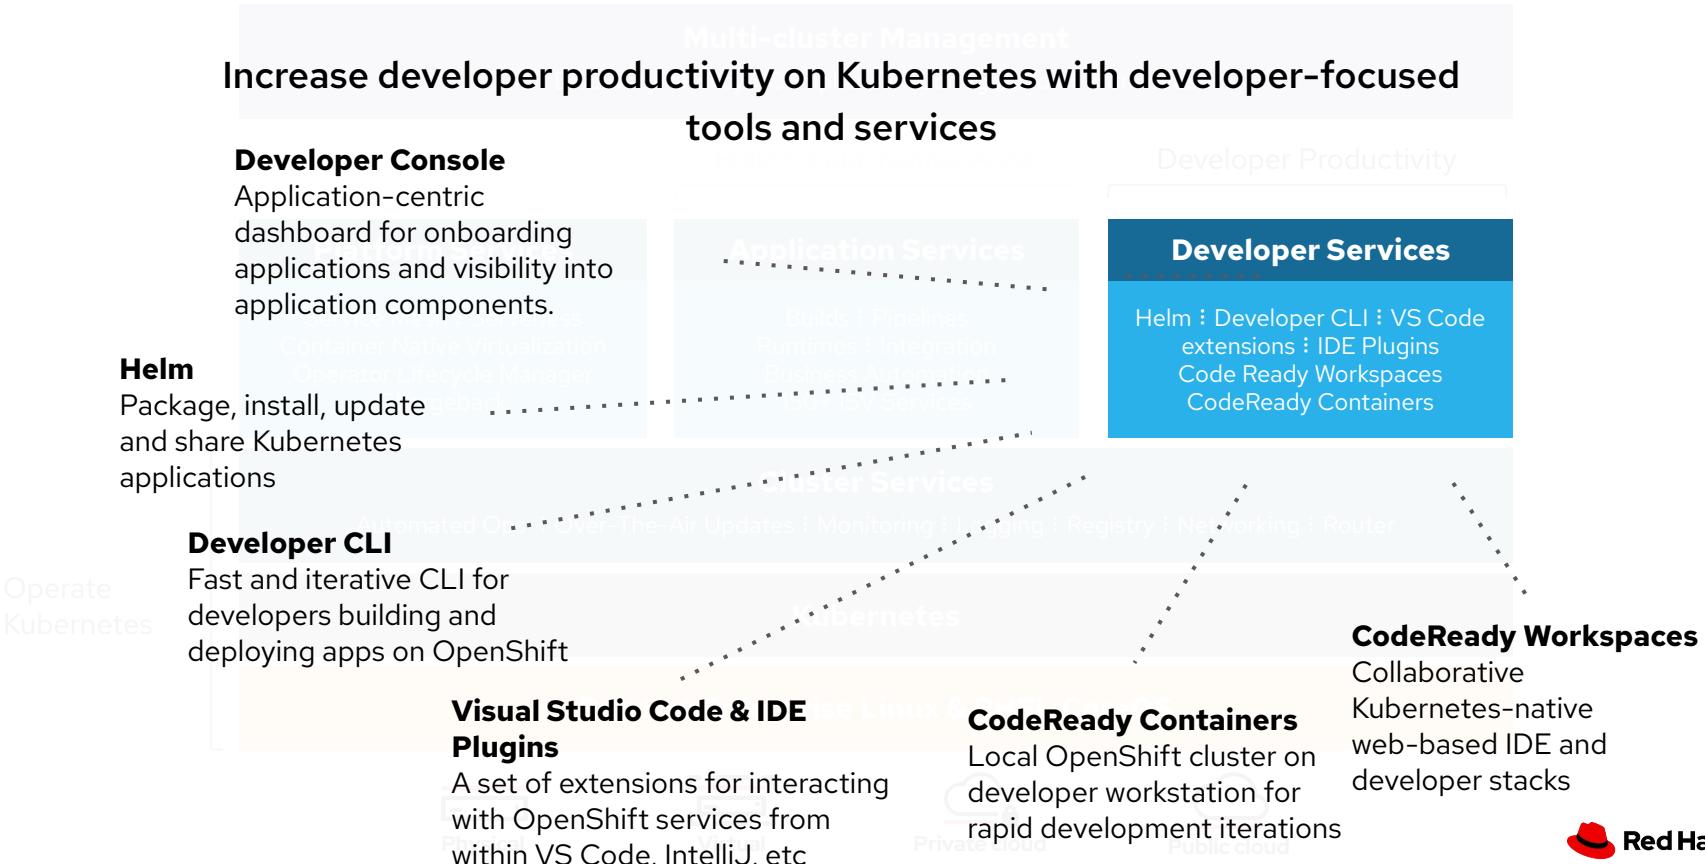


PIPELINE

- Pushing code to the Git repository triggers one or multiple deployment pipelines
- Design your pipelines based on your development workflow e.g. test the pull request
- Failure in the pipeline? Go back to the code and start again

OpenShift Developer Ecosystem

OpenShift Container Platform



Red Hat Runtimes

- Quarkus GA - Supersonic, Subatomic Java.
Native support coming in Q3 via Mandrel.
- Red Hat SSO - Support for OpenShift secret Vaults, WebAuthn protocol.
- Data Grid 8.1 - Cross-site cluster support
- Z-Series support for Runtimes - Use existing Z & Cloud Pak investment with Runtimes & OpenShift.
- JBOSS EAP expansion pack for MicroProfile
- Spring Boot 2.2 - New AMQ Starters, GA of Reactive support and Kubernetes Java annotations.

[What's New in Red Hat Runtimes \[Q2\]](#)

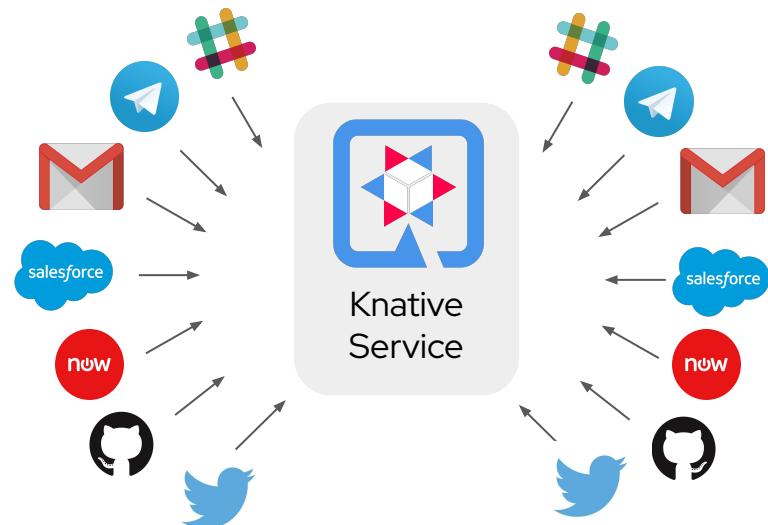
LAUNCH SERVICE

CLOUD-NATIVE RUNTIMES

LAUNCH SERVICE				
CLOUD-NATIVE RUNTIMES				

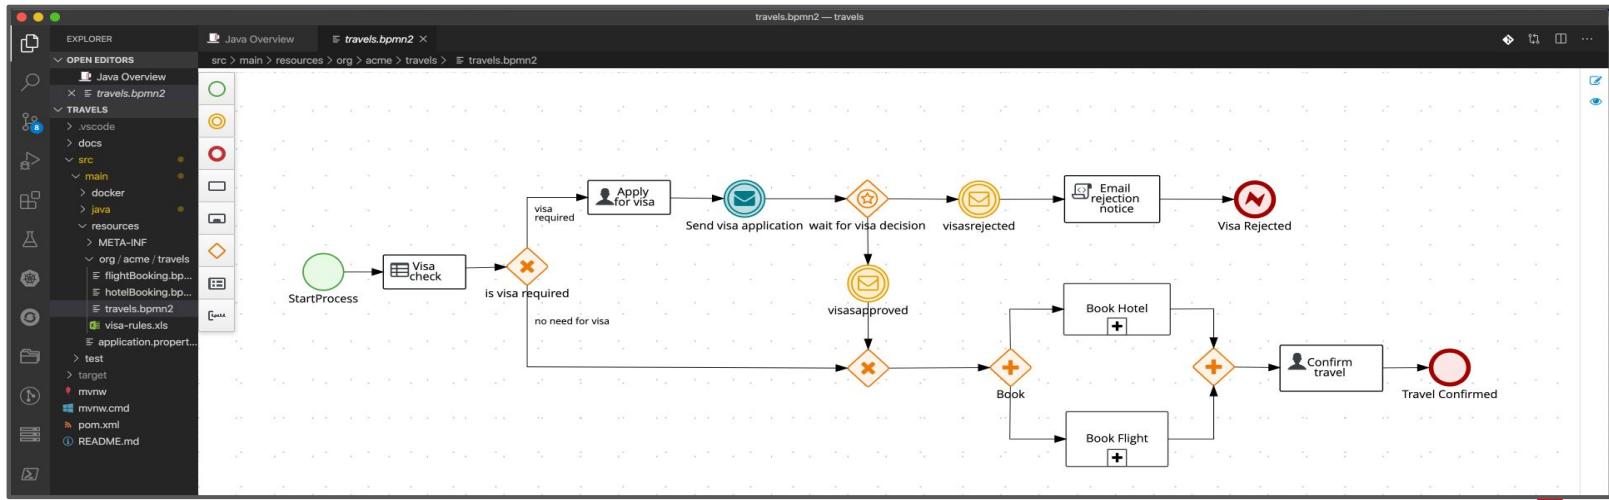
Red Hat Integration

- **Service Registry GA** - schema and API registry to support Kafka and API-based workloads on OpenShift.
- **Mirror Maker 2.0** - major evolution in data replication for Kafka on OpenShift now at full support in Strimzi
- **Camel K for Serverless (TP)** - leverage the huge Camel connector catalog to drive events into your OpenShift Serverless applications based on Camel K and Knative Eventing.



Red Hat Process Automation

- Kogito - moves to developer preview for next-gen business automation based on Quarkus
- Trusty AI - first MVP of Trusty AI in the next Kogito developer preview release. With runtime dashboards and metrics.
- Red Hat Business Automation Bundle - released in the [Visual Studio Code Marketplace](#) as Developer Preview



Quarkus Business Value



QUARKUS

Supersonic. Subatomic. Java.

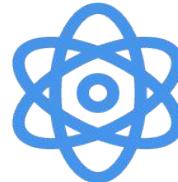
Quarkus - Kubernetes Native Java



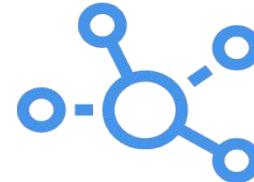
Monolith



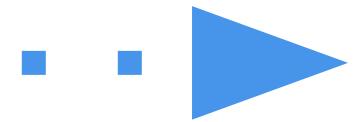
Cloud Native



Microservices



Serverless



Event-Driven
Architecture



kubernetes



Istio



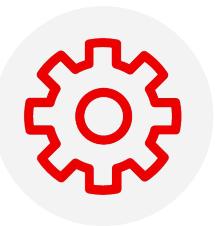
Knative

WHAT IS QUARKUS?

QUARK: elementary particle / **US**: hardest thing in computer science

Quarkus Business Value

"Supersonic, Subatomic Java"



Cost Savings

Cloud efficiency (low memory, fast startup, high density), serverless deployments



Faster Time to Market

Developer productivity, extensions ecosystem, low learning curve, keep competitive edge



Reliability

Trusted technology, active community, trusted sponsor, proven Java libraries and standards

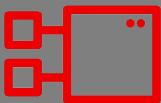
Why Quarkus?

"Quarkus is a great fit if you are looking for..."



IMPROVED APP PERFORMANCE & COST

- Reduce memory consumption
- Increase startup speeds
- Reduce application size
- Looking for Spring or Java alternatives



DIGITAL TRANSFORMATION

- Modernize or optimize existing applications
- Move to a cloud-native or microservices architecture
- Need lightweight connectivity between services
- Next generation business automation



SERVERLESS

- Have seldom used but critical services
- Interested in serverless or FaaS (function as a service)

Quarkus Technical Value

“Historical” Enterprise Java Stack

Architecture: **Monoliths**



Deployment: **multi-app, appserver**

Dynamic Application Frameworks

App Lifecycle: **Months**

Application Server

Memory: **1GB+ RAM**

Java Virtual Machine (Hotspot)

Startup Time: **10s of sec**

Operating System + Hardware/VM

“Modern” Enterprise Java Stack

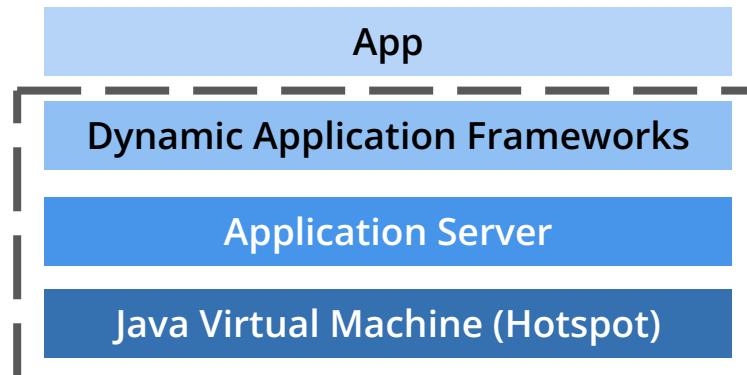
Architecture: **Microservices**

Deployment: **Single App**

App Lifecycle: **Days**

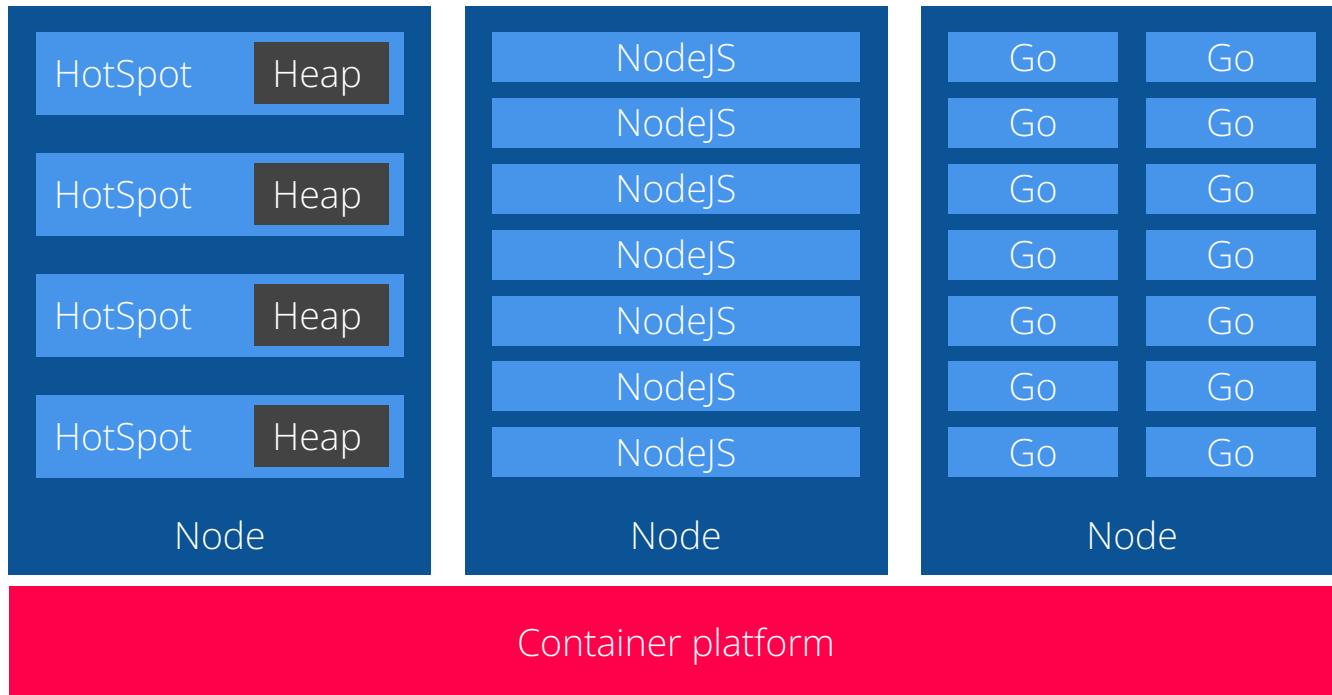
Memory: **100MBs+ RAM**

Startup Time: **Seconds**



No
Change

Hidden Truth About Java + Containers



**THERE IS A NEED FOR A
NEW JAVA STACK FOR
CLOUD-NATIVE AND
SERVERLESS**

Experts from cloud-native Java OS projects

VERT.X



Eclipse Vert.x



Hibernate



RESTEasy



Eclipse MicroProfile



Undertow

OpenJDK™



QUARKUS

Differentiators



Container First

- Tailors your app for HotSpot & GraalVM
- Fast boot time and low RSS memory
- Serverless fit



Developer Joy

- Live coding
- Unified configuration



Unifies Imperative & Reactive

- Combines blocking and non-blocking
- Built-in event bus



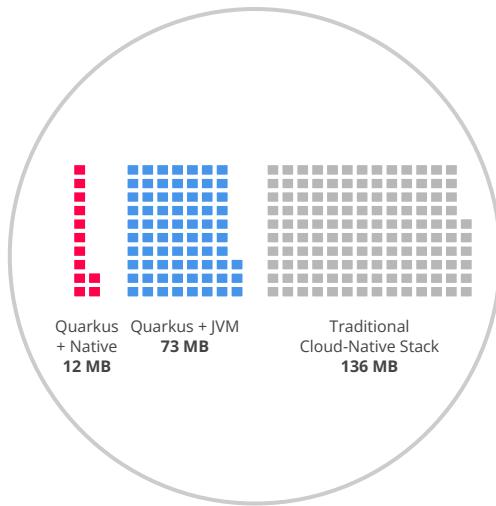
Best of Breed Libraries & Standards

- 90+ extensions
- "Powered by Quarkus" applications

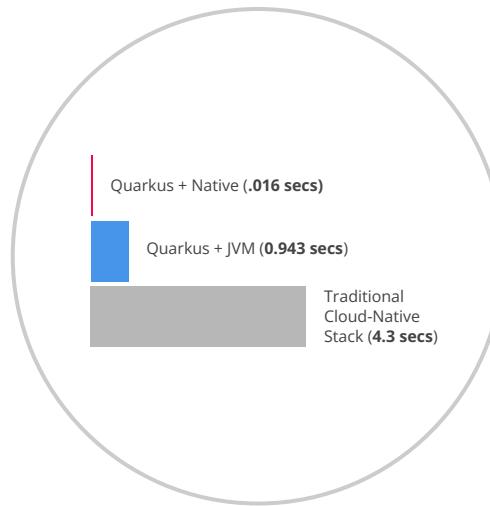
Benefit No. 1: Container First

"We went from 1-min startup times to 400 milliseconds"

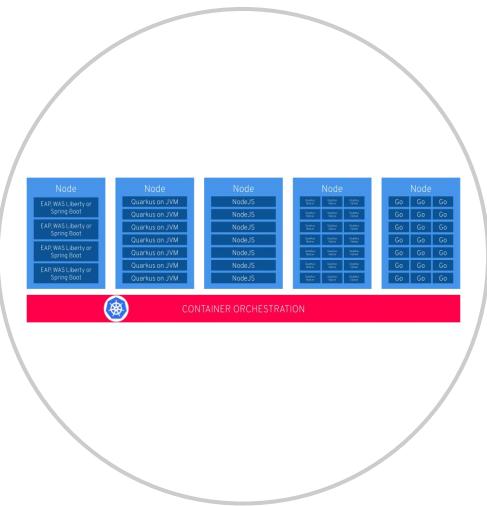
Reduced Memory Footprint



Fast Startup Time



Smaller Disk Footprint

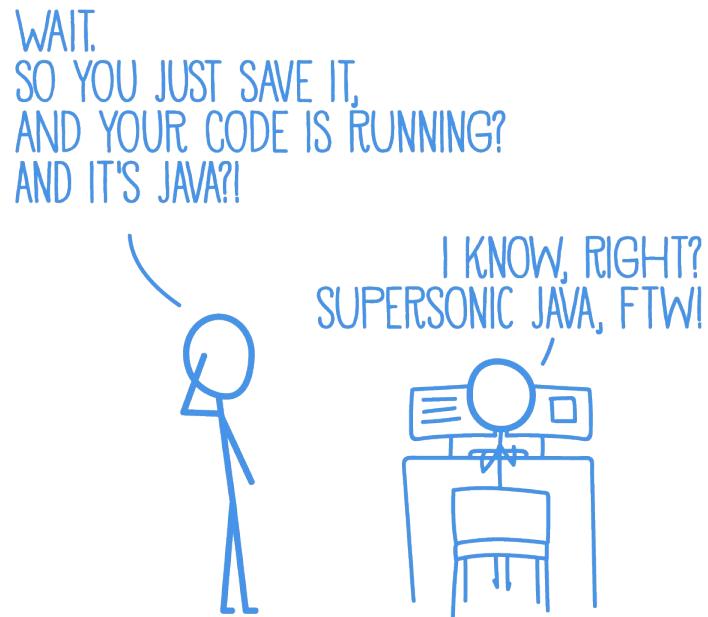


Benefit No. 2: Developer Joy

*"Our developers used to wait **2 to 3 mins** to see their changes. **Live coding** does away with this."*

A cohesive platform for optimized developer joy:

- Based on standards and more
- Unified configuration
- Live coding
- Streamlined code for the 80% common usages, flexible for the 20%
- No hassle native executable generation



Benefit No. 3: Unifies Imperative and Reactive

```
@Inject  
SayService say;  
  
@GET  
@Produces(MediaType.TEXT_PLAIN)  
public String hello() {  
    return say.hello();  
}
```

```
@Inject @Stream("kafka")  
Publisher<String> reactiveSay;  
  
@GET  
@Produces(MediaType.SERVER_SENT_EVENTS)  
public Publisher<String> stream() {  
    return reactiveSay;  
}
```

- Combine both Reactive and imperative development in the same application
- Inject the EventBus or the Vertx context
- Use the technology that fits your use-case
- Key for reactive systems based on event driven apps

Benefit No. 4: Best of Breed Frameworks & Standards

"When you adopt Quarkus, you will be productive from day one since you don't need to learn new technologies."



Eclipse Vert.x



Hibernate



RESTEasy



Apache Camel



Eclipse MicroProfile



Netty



Kubernetes



OpenShift



Jaeger



Prometheus



Apache Kafka



Infinispan



Flyway



Neo4j



MongoDB



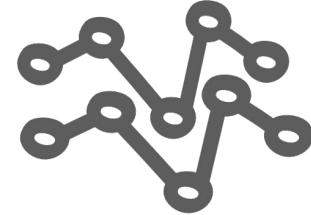
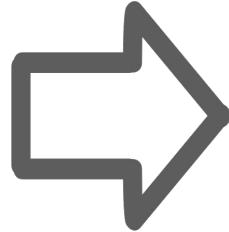
MQTT



Keycloak



Apache Tika



Supersonic, Subatomic

Fast.

Blazing fast to start.

Millisecond fast!

Supersonic, Subatomic Java

REST

Quarkus + Native (via GraalVM) **0.016 Seconds**

Quarkus + JVM (via OpenJDK) **0.943 Seconds**

Traditional Cloud-Native Stack **4.3 Seconds**

REST + CRUD

Quarkus + Native (via GraalVM) **0.042 Seconds**

Quarkus + JVM (via OpenJDK) **2.033 Seconds**

Traditional Cloud-Native Stack **9.5 Seconds**

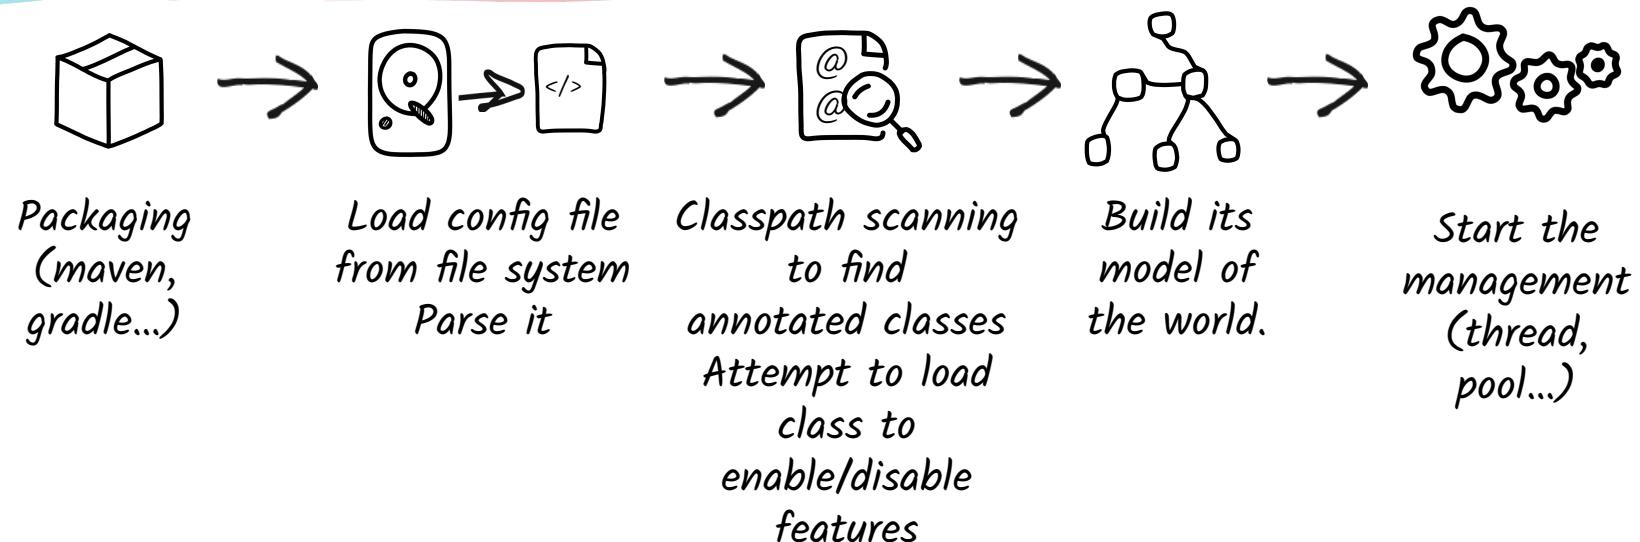
Time to first response

HOW DOES QUARKUS WORK?

How Does a Framework Start?

Build Time

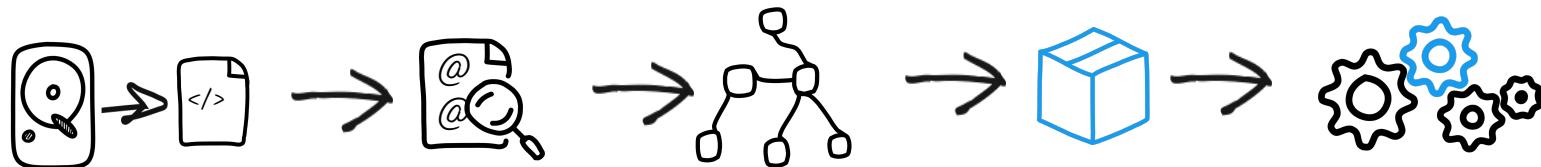
Runtime



The Quarkus Way

Build Time

Runtime

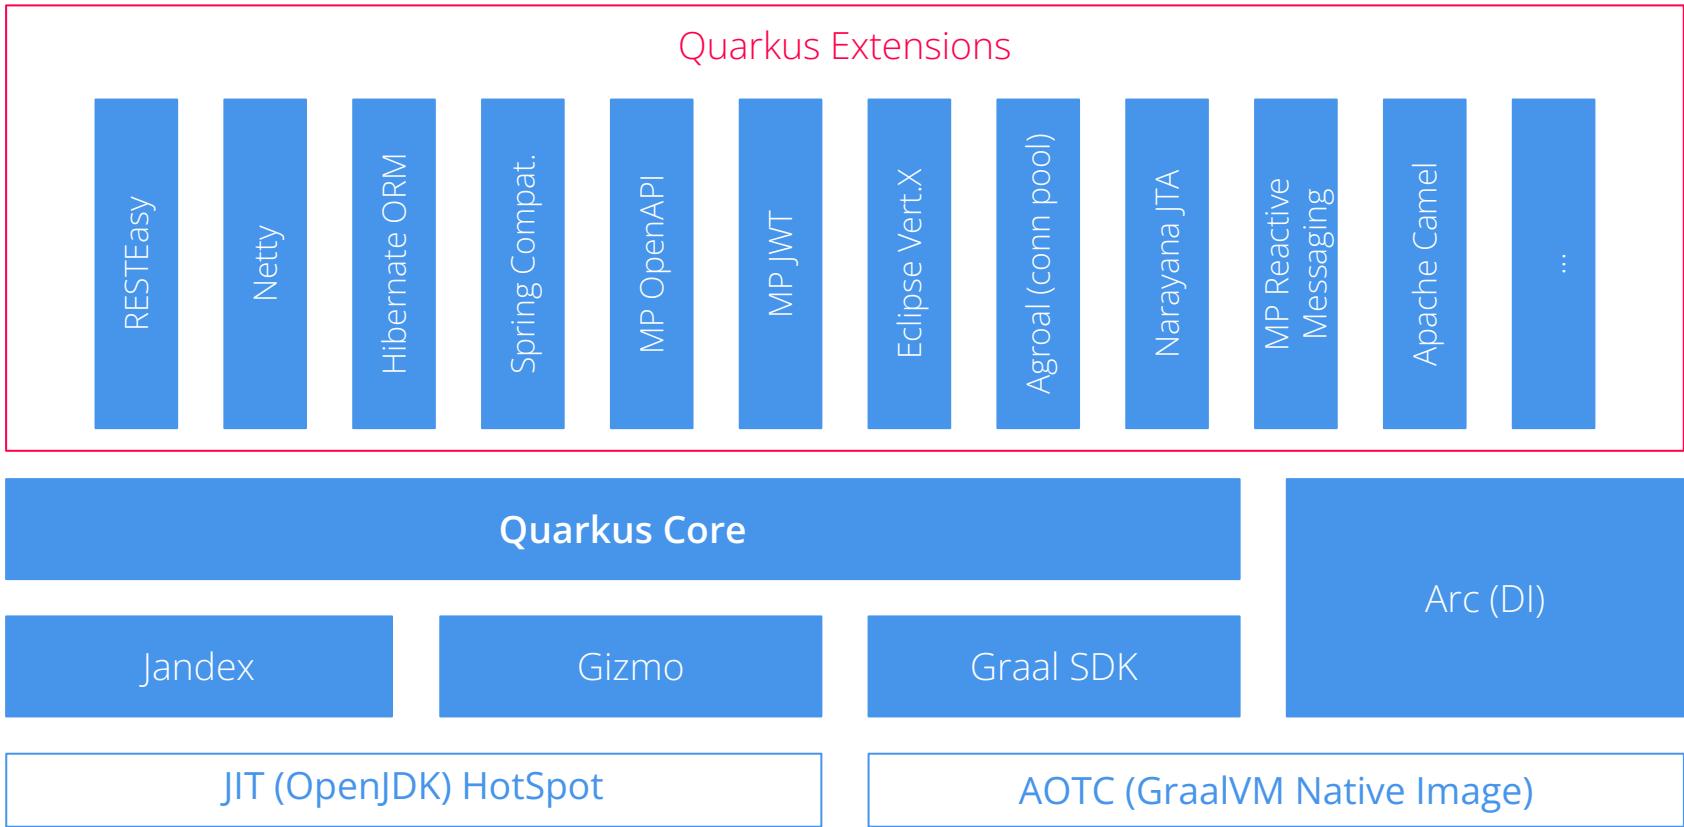


Build Time

Runtime

An ahead-of-time, build-time, runtime





The Right VM For the Right Deployment

JIT (OpenJDK HotSpot)

- High memory density requirements
- High request/s/MB
- Fast startup time
- Best raw performance (CPU)
- Best garbage collectors
- Higher heap size usage
- Known monitoring tools
- Compile Once, Run anywhere
- Libraries that only work in standard JDK

AOT (GraalVM native image)*

- Highest memory density requirements
- Highest request/s/MB
for low heap size usages
- Faster startup time
10s of ms for Serverless

*Currently in Tech Preview



Quarkus Tools - Build

maven



Gradle*

```
mvn io.quarkus:quarkus-maven-plugin:1.3.2.Final-redhat-00001:create \
-DprojectGroupId=org.acme \
-DprojectArtifactId=getting-started \
-DplatformGroupId=com.redhat.quarkus \
-DplatformVersion=1.3.2.Final-redhat-00001 \
-DclassName="org.acme.quickstart.GreetingResource" \
-Dpath="/hello"
cd getting-started
```

*community supported



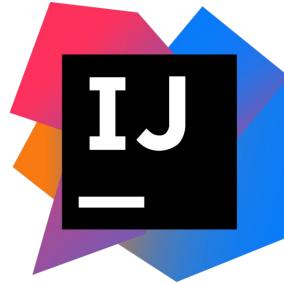
Quarkus Tools - IDE



[VSCode](#)



[Eclipse](#)



[IntelliJ](#)



[che.openshift.io](#)

DEMO

THANK YOU